

Experiment report: Data Visualization

Done by student: **Bylkova Kristina (伯汀娜), 1820249049**

Course: **Big Data Analysis Technology**

Date: **October 2024**

Task 1

Types of data analysis and what kind of Data Visualization tools can support which type of data analysis?

1. Descriptive analysis: deals with what happened in the Past (Microsoft Power BI, Tableau, Matplotlib & Seaborn);
2. Diagnostic analysis: deals with why it happened in the Past (Tableau, Power BI, D3.js, Matplotlib & Seaborn);
3. Predictive analysis: deals with what will happen in the Future (Tableau, Power BI, Plotly, Matplotlib & Seaborn);
4. Perspective analysis: deals with what step to take to avoid a problem in the future (Power BI, Tableau, D3.js, Plotly).

Task 2

Data Visualization tools:

1. Microsoft Power BI: provides real-time data insights through dashboards;
2. Tableau: a tool for creating interactive data visualizations;
3. D3.js (Data-Driven Documents): an open-source library written in javascript, creates dashboards in web browsers;
4. Plotly: a graphing library used for creating interactive, web-based data visualizations, including charts, graphs, and dashboards;
5. Matplotlib & Seaborn: Matplotlib is a Python library for creating visualizations and Seaborn is a Python data visualization library built on top of Matplotlib (it also makes visualizations).

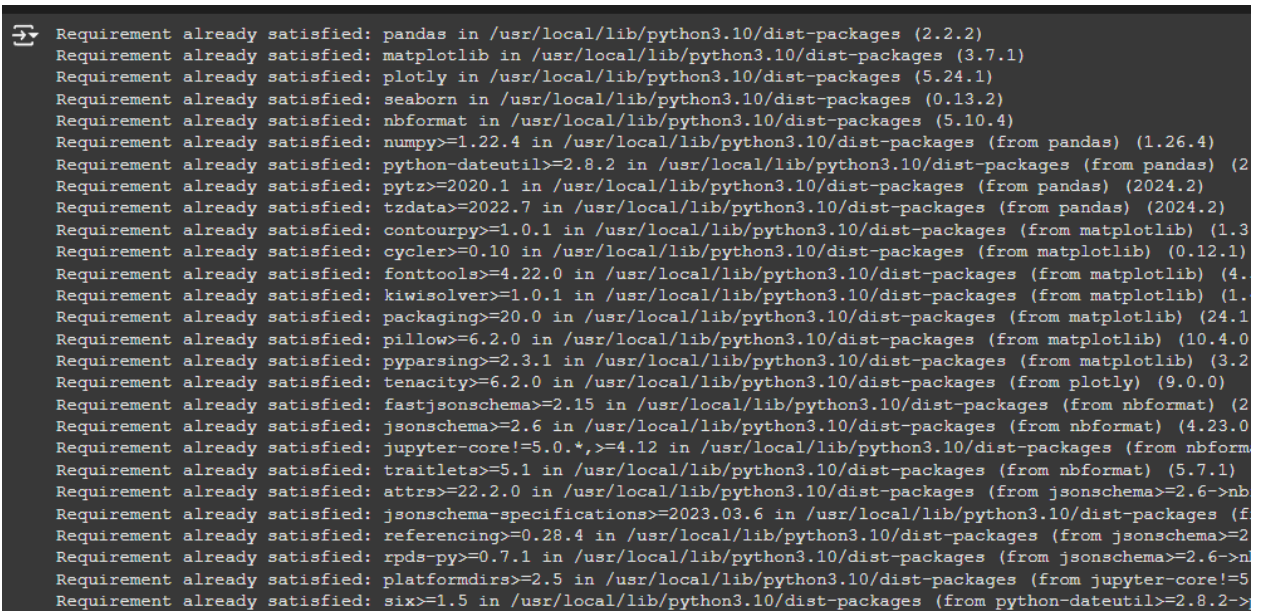
Experiment 1: Statistical Analysis

Here we dwelve in to statistical information, where we filter data through an attribute.

Code:

```
1 # Importing all necessary libraries for data manipulation, analysis and visualization
2 import sys
3 ![sys.executable] -m pip install pandas matplotlib plotly seaborn nbformat
4
5 import plotly.express as px
6 import plotly.graph_objects as go
7 import pandas as pd
8 import sqlite3
9 import matplotlib.pyplot as plt
10 import numpy as np
11 import seaborn as sns
12
13 # Creating a connection to the database and load the data
14 conn = sqlite3.connect('imdb_games.db')
15 df_cleaned = pd.read_sql_query("SELECT * FROM games", conn)
```

Result:

A terminal window with a dark background and light-colored text. It shows a list of requirements for various Python libraries, all of which are marked as 'Requirement already satisfied'. The libraries listed include pandas, matplotlib, plotly, seaborn, nbformat, numpy, python-dateutil, pytz, tzdata, contourpy, cycler, fonttools, kiwisolver, packaging, pillow, pyparsing, tenacity, fastjsonschema, jsonschema, jupyter-core, traitlets, attrs, jsonschema-specifications, referencing, rpds-py, platformdirs, and six. Each line specifies the version of the library and the source (e.g., 'from pandas' or 'from matplotlib').

```
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (2.2.2)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (3.7.1)
Requirement already satisfied: plotly in /usr/local/lib/python3.10/dist-packages (5.24.1)
Requirement already satisfied: seaborn in /usr/local/lib/python3.10/dist-packages (0.13.2)
Requirement already satisfied: nbformat in /usr/local/lib/python3.10/dist-packages (5.10.4)
Requirement already satisfied: numpy>=1.22.4 in /usr/local/lib/python3.10/dist-packages (from pandas) (1.26.4)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas) (2024.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/dist-packages (from pandas) (2024.2)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.3.0)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (4.53.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.4.5)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (24.1)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (10.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (3.2.0)
Requirement already satisfied: tenacity>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from plotly) (9.0.0)
Requirement already satisfied: fastjsonschema>=2.15 in /usr/local/lib/python3.10/dist-packages (from nbformat) (2.20.1)
Requirement already satisfied: jsonschema>=2.6 in /usr/local/lib/python3.10/dist-packages (from nbformat) (4.23.0)
Requirement already satisfied: jupyter-core!=5.0.*,>=4.12 in /usr/local/lib/python3.10/dist-packages (from nbformat) (5.7.1)
Requirement already satisfied: traitlets>=5.1 in /usr/local/lib/python3.10/dist-packages (from nbformat) (5.7.1)
Requirement already satisfied: attrs>=22.2.0 in /usr/local/lib/python3.10/dist-packages (from jsonschema) (24.2.0)
Requirement already satisfied: jsonschema-specifications>=2023.03.6 in /usr/local/lib/python3.10/dist-packages (from jsonschema) (2024.10.1)
Requirement already satisfied: referencing>=0.28.4 in /usr/local/lib/python3.10/dist-packages (from jsonschema) (0.35.0)
Requirement already satisfied: rpds-py>=0.7.1 in /usr/local/lib/python3.10/dist-packages (from jsonschema) (0.20.1)
Requirement already satisfied: platformdirs>=2.5 in /usr/local/lib/python3.10/dist-packages (from jupyter-core) (4.2.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil) (1.16.0)
```

Figure 1: Importing libraries

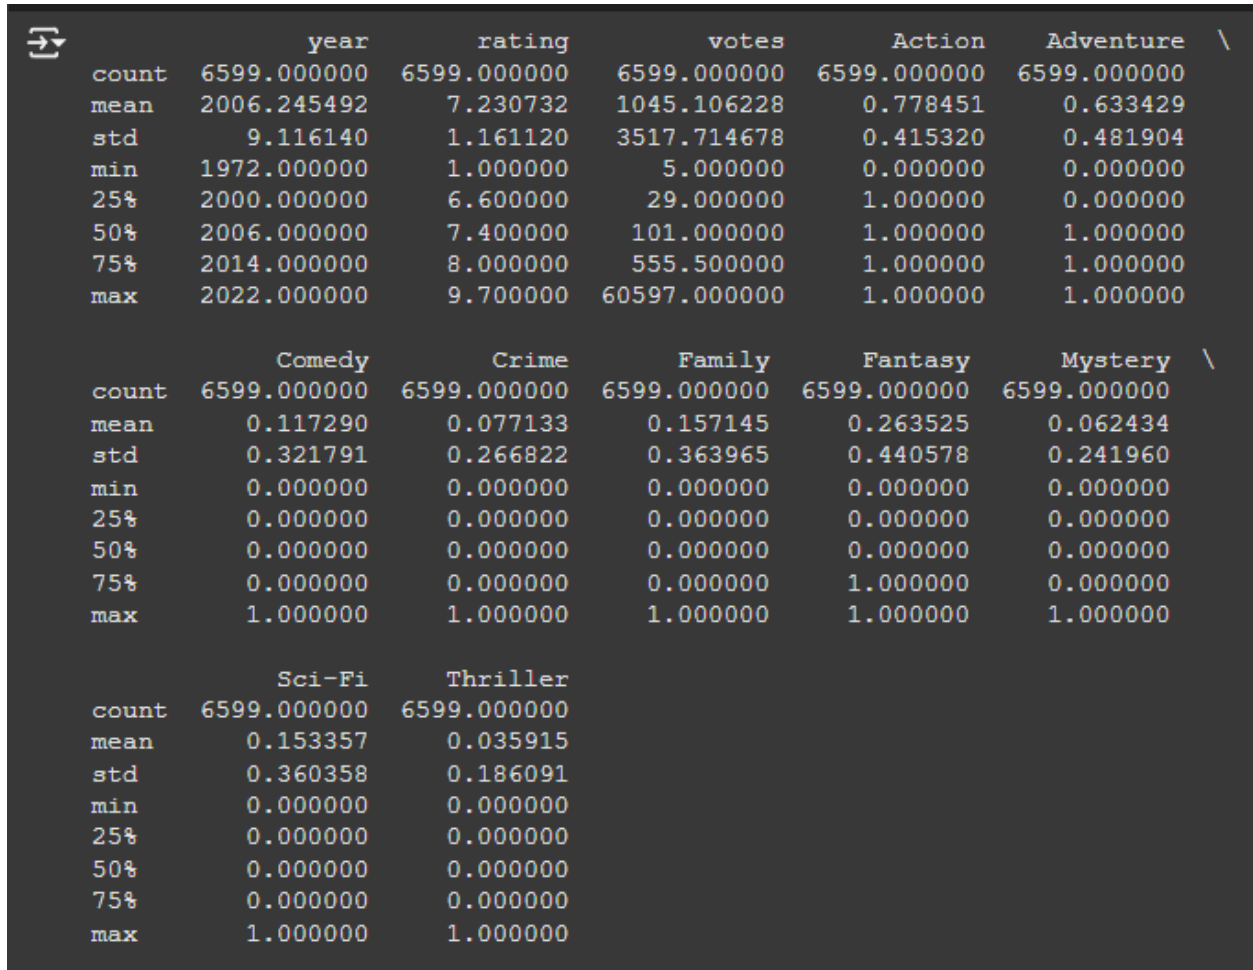
Experiment 1.1: Descriptive Statistics

Getting the descriptive statistics, that is the Count, Mean, Standard Deviation, Percentiles (25%, 50%, 75%), minimum and maximum values of each column of the cleaned data from the dataset.

Code:

```
1 summary_stats = df_cleaned.describe()
2 print(summary_stats)
```

Result:



The image shows a terminal window with a dark background and light-colored text. It displays the output of the pandas `describe()` method applied to a DataFrame named `df_cleaned`. The output is organized into three sections, each starting with a genre header. The first section covers the 'year' column and the 'Action' and 'Adventure' genres. The second section covers the 'Comedy', 'Crime', 'Family', 'Fantasy', and 'Mystery' genres. The third section covers the 'Sci-Fi' and 'Thriller' genres. Each section lists eight statistics: count, mean, std, min, 25%, 50%, 75%, and max. The 'year' column has a count of 6599 and a mean of approximately 2006.25. The 'Action' and 'Adventure' genres have counts of 6599 and means around 0.78 and 0.63 respectively. The other genres have counts of 6599 and means close to 0. The 'Sci-Fi' and 'Thriller' genres have counts of 6599 and means around 0.15 and 0.04 respectively.

| | year | rating | votes | Action | Adventure | \ |
|-------|-------------|-------------|--------------|-------------|-------------|---|
| count | 6599.000000 | 6599.000000 | 6599.000000 | 6599.000000 | 6599.000000 | |
| mean | 2006.245492 | 7.230732 | 1045.106228 | 0.778451 | 0.633429 | |
| std | 9.116140 | 1.161120 | 3517.714678 | 0.415320 | 0.481904 | |
| min | 1972.000000 | 1.000000 | 5.000000 | 0.000000 | 0.000000 | |
| 25% | 2000.000000 | 6.600000 | 29.000000 | 1.000000 | 0.000000 | |
| 50% | 2006.000000 | 7.400000 | 101.000000 | 1.000000 | 1.000000 | |
| 75% | 2014.000000 | 8.000000 | 555.500000 | 1.000000 | 1.000000 | |
| max | 2022.000000 | 9.700000 | 60597.000000 | 1.000000 | 1.000000 | |

| | Comedy | Crime | Family | Fantasy | Mystery | \ |
|-------|-------------|-------------|-------------|-------------|-------------|---|
| count | 6599.000000 | 6599.000000 | 6599.000000 | 6599.000000 | 6599.000000 | |
| mean | 0.117290 | 0.077133 | 0.157145 | 0.263525 | 0.062434 | |
| std | 0.321791 | 0.266822 | 0.363965 | 0.440578 | 0.241960 | |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | |
| 25% | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | |
| 50% | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | |
| 75% | 0.000000 | 0.000000 | 0.000000 | 1.000000 | 0.000000 | |
| max | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | |

| | Sci-Fi | Thriller |
|-------|-------------|-------------|
| count | 6599.000000 | 6599.000000 |
| mean | 0.153357 | 0.035915 |
| std | 0.360358 | 0.186091 |
| min | 0.000000 | 0.000000 |
| 25% | 0.000000 | 0.000000 |
| 50% | 0.000000 | 0.000000 |
| 75% | 0.000000 | 0.000000 |
| max | 1.000000 | 1.000000 |

Figure 2: Descriptive Statistics


Experiment 1.2: Top 11 games by ratings

Discovering the top 11 games by ratings from the cleaned dataset.

Code:

```
1 conn = sqlite3.connect('imdb_games.db')
2
3 query = '''
4 SELECT name, rating
5 FROM games
6 ORDER BY rating DESC
7 LIMIT 11
8 '''
9 top_games = pd.read_sql(query, conn)
10
11 conn.close()
12
13 print(top_games.to_string(index=False))
```

Result:



| name | rating |
|---|--------|
| Red Dead Redemption II | 9.7 |
| The Last of Us | 9.7 |
| The Witcher 3: Wild Hunt | 9.7 |
| Mass Effect: Legendary Edition | 9.7 |
| The Witcher 3: Wild Hunt - Blood and Wine | 9.7 |
| The Last of Us | 9.7 |
| The Witcher 3: Wild Hunt - Blood and Wine | 9.7 |
| Red Dead Redemption II | 9.7 |
| The Witcher 3: Wild Hunt | 9.7 |
| The Last of Us | 9.7 |
| Mass Effect: Legendary Edition | 9.7 |

Figure 3: Top 11 games by ratings


Experiment 1.3: Total Votes for Each Game

Discovering the total votes for each game genre.

Code:

```
1 # Defining the genre columns
2 genre_columns = ['Action', 'Adventure', 'Comedy', 'Crime', 'Family',
3                 'Fantasy', 'Mystery', 'Sci-Fi', 'Thriller']
4
5 melted_df = df_cleaned.melt(id_vars=['votes'], value_vars=genre_columns,
6                             var_name='Genre', value_name='Is_Genre')
7
8 # Filtering only applicable genres
9 filtered_df = melted_df[melted_df['Is_Genre'] == 1]
10
11 # Group by Genre and sum votes
12 total_votes_by_genre = filtered_df.groupby('Genre')['votes'].sum().reset_index()
13
14 # Sort by total votes in descending order
15 total_votes_by_genre = total_votes_by_genre.sort_values(by='votes', ascending=False)
16
17 print(total_votes_by_genre)
```

Result:



| | Genre | votes |
|---|-----------|---------|
| 0 | Action | 6348811 |
| 1 | Adventure | 4977777 |
| 5 | Fantasy | 1631520 |
| 3 | Crime | 1343081 |
| 7 | Sci-Fi | 982141 |
| 2 | Comedy | 457004 |
| 8 | Thriller | 440125 |
| 6 | Mystery | 429245 |
| 4 | Family | 368054 |

Figure 4: Total votes for each genre

Experiment 2: Visualization

Experiment 2.1: Distribution of ratings

Getting the distribution of ratings with the use of a histogram.

Code:

```
1 mean_rating = df_cleaned['rating'].mean()
2
3 fig = px.histogram(df_cleaned, x='rating', nbins=90,
4                   title='Distribution of Video Game Ratings',
5                   labels={'rating': 'Rating'},
6                   color_discrete_sequence=['purple'])
7
8 fig.add_vline(x=mean_rating, line_color='red', line_dash='dash',
9              annotation_text=f'Mean: {mean_rating:.2f}',
10             annotation_position='top right')
11
12 fig.update_layout(xaxis_title='Rating', yaxis_title='Frequency')
13
14 fig.show()
```

Result:




Figure 5: Distribution of ratings

Experiment 2.2: Ratings over the years with trendline

Getting the ratings over the years using a line plot aided with a trendline.

Code:

```
1 average_ratings = df_cleaned.groupby('year')['rating'].mean().reset_index()
2
3 fig = px.line(average_ratings, x='year', y='rating',
4               title='Average Ratings of Video Games Over the Years',
5               labels={'year': 'Year', 'rating': 'Average Rating'},
6               markers=True)
7
8 trendline = go.Scatter(x=average_ratings['year'], y=average_ratings['rating'],
9                        mode='lines',
10                       name='Trendline',
11                       line=dict(color='red'))
12
13 fig.add_trace(trendline)
14
15 fig.show()
```

Result:




Figure 6: Ratings over the years with trendline

Experiment 3: Optional Experiment

Experiment 3.1: How are certificates related to different genres?

Exploring the relationship between game certificates and all the genres using a grouped bar chart to visualize.

Code:

```
1 melted_df = df_cleaned.melt(id_vars=['certificate', 'rating'], value_vars=genre_columns,
2                             var_name='Genre', value_name='Is_Genre')
3
4 filtered_df = melted_df[melted_df['Is_Genre'] == 1]
5
6 # Grouping by certificate and genre to calculate the average rating
7 cert_genre_relation = filtered_df.groupby(['certificate', 'Genre'])['rating'].mean().reset_index()
8
9 fig = px.bar(cert_genre_relation, x='certificate', y='rating', color='Genre',
10             title='Average Genre Ratings by Certificate',
11             labels={'certificate': 'Certificate', 'rating': 'Average Rating'},
12             barmode='group')
13
14 fig.update_layout(xaxis_title='Certificate',
15                  yaxis_title='Average Rating',
16                  legend_title='Genre')
17
18 fig.show()
```

Result:




Figure 7: Relation of certificates

Experiment 3.2: What is the relationship between the Genres, ratings, release year and votes?

Making a correlation heatmap.

Code:

```
1 numeric_df = df_cleaned.select_dtypes(include=[np.number])
2
3 plt.figure(figsize=(10, 8))
4 correlation_matrix = numeric_df.corr() # Calculating correlation on numeric DataFrame
5 sns.heatmap(correlation_matrix, annot=True, fmt=".2f", cmap='coolwarm')
6 plt.title('Correlation Heatmap')
7 plt.show()
```

Result:




Figure 8: Correlation Heatmap